


Introduction to Kubernetes

Remko de Knikker
Developer Advocate

What is
container
orchestration?

What is container orchestration?

- Container orchestration
- Cluster management
- Scheduling
- Service discovery
- Replication
- Health management



What is Kubernetes?

Container Orchestrator

- Provision, manage, scale applications
- Manage infrastructure resources needed by applications
 - Volumes
 - Networks
 - Secrets
 - And many, many, many more..
- What's in a name?
 - Kubernetes (K8s/Kube): "Helmsman" in ancient Greek
- Declarative model
 - Provide the "desired state" and Kubernetes will make it happen




How was Kubernetes created

- Based on Google's Borg & Omega
- Open Governance
 - Cloud Native Compute Foundation
- Adoption by Enterprise
 - RedHat, Microsoft, IBM and Amazon

Kubernetes Architecture

- At its core, Kubernetes is a database (etcd), with "watchers" & "controllers" that react to changes in the DB
- The controllers are what make it Kubernetes
- This pluggability and extensibility is part of its "secret sauce".
- DB represents the user's desired state
 - Watchers attempt to make reality match the desired state



Kubernetes Resource Model

- Config Maps
 - Daemon Sets
 - **Deployments**
 - Events
 - Endpoints
 - Ingress
 - Jobs
 - Nodes
 - Namespaces
 - **Pods**
 - Persistent Volumes
 - Replica Sets
 - Secrets
 - Service Accounts
 - **Services**
 - Stateful Sets, and more...
- Kubernetes aims to have the building blocks on which you build a cloud native platform.
 - Therefore, the internal resource model **is** the same as the end user resource model.
- ### Key Resources
- Pod: set of co-located containers
 - Smallest unit of deployment
 - Several types of resources to help manage them
 - Replica Sets, Deployments, Stateful Sets, ...
 - Services
 - Define how to expose your app as a DNS entry
 - Query based selector to choose which pods apply

Kubernetes Client

CLI tool to interact with Kubernetes cluster

Platform specific binary available to download


- <https://kubernetes.io/docs/tasks/tools/install-kubectl>

The user directly manipulates resources via json/yaml

```
$ kubectl (create|get|apply|delete) -f myResource.yaml
```


Kubernetes in Action

1. User via "kubectl" deploys a new application
2. API server receives the request and stores it in the DB (etcd)
3. Watchers/controllers detect the resource changes and act upon it
4. ReplicaSet watcher/controller detects the new app and creates new pods to match the desired # of instances
5. Scheduler assigns new pods to a kubelet
6. Kubelet detects pods and deploys them via the container running (e.g. Docker)
7. KubeProxy manages network traffic for the pods – including service discovery and load-balancing



Benefits of Container Orchestration

- Automated scheduling and scaling
- Zero downtime deployments
- High availability and fault tolerance
- A/B deployments



But Wait? What About Production?

- Kubernetes by itself is not enterprise-ready
- Kubernetes must integrate with underlying platform to provide infrastructure, storage, etc.
- Lacking in operational view + controls, pre-built catalogs

Lab Time