


# Intro to Kafka and Event Streams

Remko de Knikker  
Developer Advocate

# IBM Watson and cloud Platform



IBM Event Streams is built on Apache Kafka



“Kafka adoption was accelerating...with production deployments of Kafka claiming **six of the top 10** travel companies, **seven of the top 10** global banks, **eight of the top 10** insurance companies, and **nine of the top 10** US telecom companies.”

# IBM Event Streams is built on Apache Kafka


Adding years of operational expertise of running Apache Kafka for enterprises.

- Enterprise Connectors
- Geo-replication
- Enterprise grade security
- Schema registry with support for Apache Avro schemas
- Tight Integration with IBM IoT Platform, Object Storage, Streaming Analytics, Cloud Functions
- Scalable REST API
- CLI support for automation
- Enterprise Platform support: Red Hat OpenShift, IBM Z
- Fully Managed or Container certified,
- Zero-down time upgrade
- 24/7 Support

# Apache Kafka

# High-throughput message bus

# Kafka is a Streaming Platform



# Apache Kafka

A streaming platform has three key capabilities:

- Publish and subscribe to streams of data, similar to a message queue or enterprise messaging system.
- Store streams of data in a fault-tolerant durable way.
- Process streams of data on-demand.

# Apache Kafka

Kafka is generally used for:

- Real-time, reliable streaming data pipelines between systems
- Real-time streaming applications that transform or react to the streams of data

# Apache Kafka

Kafka includes five core APIs:


- The [Producer](#) API to send streams of data to topics.
- The [Consumer](#) API to read streams of data from topics.
- The [Streams](#) API allows transforming streams of data from input to output topics.
- The [Connect](#) API for implementing connectors to pull from a source system or push into a sink system.
- The [AdminClient](#) API for managing and inspecting topics, brokers, and other Kafka objects.

# Use Cases

## Worker Offload

Intensive work offloaded and distributed amongst worker processes to be performed asynchronously


- Processing images or videos
- Performing text analytics



## Event-Driven

Take one or more actions when something interesting happens

- Email logs and update dashboards when build finishes
- Upload videos once finished transcoding




# Use Cases

## Delayed Processing

Schedule a task to happen at a specific time


- Run detailed reports when app use is low
- Generate end of day summary



## 3<sup>rd</sup> Party Integration


Ensure applications remain responsive even when 3<sup>rd</sup> party systems are not available or responding fast enough

- Updating existing CRM system
- Booking appointment




# Kafka is: Built for Scale

Each partition is an ordered, immutable sequence of records that is continually appended to a structured commit log.




# Kafka is: Built for Scale

Records in the partitions are each assigned a sequential id number called the *offset*.



# Kafka partitions messages for scale



Publishers can influence partitioning.

Each partition is an ordered list of messages.


Partitioning and ordering are linked concepts in Kafka.

Each topic consists of 1 or more partitions.


This topic has 3 partitions so the load can be distributed across 3 servers.

Partitions allow the log to scale beyond the size of a single server. Each individual partition must fit on the hosting server.

# Queues vs topics




## Publish-Subscribe and Queueing



- Each message has 1 consumer
- Gives easy sharing among consumers


- Each message has n consumers
- Hybrid models such as shared subscriptions or consumer groups exist

# Workload distribution in Kafka




# Kafka is: Built for Scale

Kafka is run as a cluster with one or more nodes or servers.




# Kafka is: Built for Scale

Kafka is run as a cluster with one or more nodes or servers, each node with one or more partitions.




# Kafka is: Built for Scale

Kafka is run as a cluster with one or more nodes or servers, each node with one or more partitions. A cluster can span multiple datacenters.



# Kafka replicates across the nodes in a cluster




Each partition has one server acting as the "leader" and zero or more servers as "followers".

This is synchronous replication of ordered data streams. It is not just eventually consistent.

Each partition is replicated across a configurable number of servers for fault tolerance.

Kafka MirrorMaker provides geo-replication support for your clusters.


# When a node fails, leaders are reassigned



Each partition has a leader responsible for replicating the messages, in order, to the other replicas.

Recovery from node failure is very fast. There is no data loss and no need to wait for propagation of changes to get consistency restored.

# IBM Event Streams





# Catalog

[All Categories](#) **kafka**

Filter

[DevOps](#)[Operations](#)[Security](#)[Network](#)[Data Science & Analytics](#)[Data](#)[Integration](#)[Storage](#)[Blockchain](#)**ibm-eventstreams-dev** IBM Event Streams based on Apache  
Kafka.

ibm-charts

[Getting started](#)[Topics](#)[Consumer groups](#)[Monitor](#)[Toolbox](#)[Connect to this cluster](#) 🔍

# Welcome to IBM Event Streams, let's get you up and running...



## Use a simulated topic

Start exploring what IBM Event Streams has to offer with our simulated topic. You can do this even if your brokers aren't ready




## Generate a starter application

Download and install our starter Kafka application and view data flowing to and from IBM Event Streams in just a few minutes

[Learn more...](#)[FAQs](#) [GitHub](#) [Documentation](#)

### Kafka basics

Learn the basics of Apache Kafka, the heart of Event Streams.



System is healthy

# Kafka Clients

## and APIs

# Kafka APIs

## Core APIs

### – **Producer**

- send messages to a topic

### – **Consumer**

- get messages from a topic

### – **Streams**

- transform streams from input topics to output topics

### – **Connect**

- pull from a source system / application into Kafka or push to a sink system /application

### – **Admin Client APIs**

## REST Proxy

<https://github.com/confluentinc/kafka-rest>

Other language clients (external to kafka.apache.org )

- librdkafka ( C/C++)
- Node.js
- PHP
- Python
- Ruby
- C#/ .NET
- Go
- Rust
- Haskell
- OCaml

# Producer API using Kafka Console Tools

## Kafka.properties

```
sasl.jaas.config=org.apache.kafka.common.security.plain.PlainLoginModule required username="USER" password="PASSWORD";  
security.protocol=SASL_SSL  
sasl.mechanism=PLAIN  
ssl.protocol=TLSv1.2  
ssl.enabled.protocols=TLSv1.2  
ssl.endpoint.identification.algorithm=HTTPS
```

```
$ bash kafka-console-producer.sh --broker-list broker-1-a1bc2d3efg4hijkl.kafka.svc01.us-  
south.eventstreams.cloud.ibm.com:9999,  
broker-2-a1bc2d3efg4hijkl.kafka.svc01.us-south.eventstreams.cloud.ibm.com:9999,  
broker-3-a1bc2d3efg4hijkl.kafka.svc01.us-south.eventstreams.cloud.ibm.com:9999,  
broker-4-a1bc2d3efg4hijkl.kafka.svc01.us-south.eventstreams.cloud.ibm.com:9999,  
broker-5-a1bc2d3efg4hijkl.kafka.svc01.us-south.eventstreams.cloud.ibm.com:9999,  
broker-6-a1bc2d3efg4hijkl.kafka.svc01.us-south.eventstreams.cloud.ibm.com:9999  
--producer.config mykafka.properties --topic greetings
```

# Producer API

```
import org.apache.kafka.clients.producer.KafkaProducer;  
import org.apache.kafka.clients.producer.ProducerRecord;  
import org.apache.kafka.clients.producer.RecordMetadata;  
  
public KafkaProducer<byte[], byte[]> kafkaProducer;  
Properties props = new Properties();  
  
kafkaProducer = new KafkaProducer<byte[], byte[]>(props);  
  
ProducerRecord<byte[], byte[]> record = new ProducerRecord<byte[], byte[]>(topic, key.getBytes("UTF-8"),  
    value.getBytes("UTF-8));  
  
RecordMetadata m = kafkaProducer.send(record).get();  
  
System.out.println("Message produced, offset: " + m.offset());
```

The producer is *thread safe*, sharing a single instance across threads will be faster

# Producer API using Spring Boot for Kafka

```
import org.slf4j.Logger;
import org.slf4j.LoggerFactory;
import org.springframework.beans.factory.annotation.Autowired;
import org.springframework.kafka.core.KafkaTemplate;
import org.springframework.web.bind.annotation.GetMapping;
import org.springframework.web.bind.annotation.PathVariable;
import org.springframework.web.bind.annotation.RestController;
import java.util.List;
import java.util.concurrent.CopyOnWriteArrayList;

@RestController
public class EventsStreamController {
    @Autowired
    private KafkaTemplate<String, String> template;
    private List<String> messages = new CopyOnWriteArrayList<>();
    @GetMapping(value = "/send/{msg}")
    public void send(@PathVariable String msg) throws Exception {
        this.template.sendDefault(msg);
    }
}
```

# Consumer API using Kafka Console Tools

Kafka.properties

```
sasl.jaas.config=org.apache.kafka.common.security.plain.PlainLoginModule required username="USER"  
password="PASSWORD";  
security.protocol=SASL_SSL  
sasl.mechanism=PLAIN  
ssl.protocol=TLSv1.2  
ssl.enabled.protocols=TLSv1.2  
ssl.endpoint.identification.algorithm=HTTPS
```

```
bash kafka-console-consumer.sh --bootstrap-server broker-1-a1bc2d3efg4hijkl.kafka.svc01.us-  
south.eventstreams.cloud.ibm.com:9999,  
broker-2-a1bc2d3efg4hijkl.kafka.svc01.us-south.eventstreams.cloud.ibm.com:9999,  
broker-3-a1bc2d3efg4hijkl.kafka.svc01.us-south.eventstreams.cloud.ibm.com:9999,  
broker-4-a1bc2d3efg4hijkl.kafka.svc01.us-south.eventstreams.cloud.ibm.com:9999,  
broker-5-a1bc2d3efg4hijkl.kafka.svc01.us-south.eventstreams.cloud.ibm.com:9999,  
broker-6-a1bc2d3efg4hijkl.kafka.svc01.us-south.eventstreams.cloud.ibm.com:9999  
--consumer.config ../mykafka.properties --topic greetings -group 1
```

# Consumer API

```
import org.apache.kafka.clients.consumer.KafkaProducer;
import org.apache.kafka.clients.consumer.ConsumerRecord;
import org.apache.kafka.common.serialization.ByteArrayDeserializer;

class ConsumerRunnable implements Runnable {

    private KafkaConsumer<byte[], byte[]> kafkaConsumer;
    ...

    ConsumerRunnable(...) {
        Properties props = new Properties();

        kafkaConsumer = new KafkaConsumer<byte[], byte[]>(props, new ByteArrayDeserializer(), new ByteArrayDeserializer());
        kafkaConsumer.subscribe(topicList);
    }

    run() {
        while (...) {
            ConsumerRecord<byte[], byte[]> record = kafkaConsumer.poll(1000);
            System.out.println("Message consumed: " + record.value());
        }
    }
}
```

The consumer is **not thread safe**. The simple approach is to give each thread its own instance.

# Consumer API using Spring Boot for Kafka

```
import org.apache.kafka.clients.consumer.ConsumerRecord;
import org.slf4j.Logger;
import org.slf4j.LoggerFactory;
import org.springframework.beans.factory.annotation.Autowired;
import org.springframework.kafka.annotation.KafkaListener;
import org.springframework.kafka.core.KafkaTemplate;
import org.springframework.web.bind.annotation.GetMapping;
import org.springframework.web.bind.annotation.PathVariable;
import org.springframework.web.bind.annotation.RestController;
import java.util.List;
import java.util.concurrent.CopyOnWriteArrayList;

@RestController
public class EventsStreamController {
    @Autowired
    private KafkaTemplate<String, String> template;
    private List<String> messages = new CopyOnWriteArrayList<>();
    @KafkaListener(topics = "${listener.topic}", groupId = "channel1")
    public void listen(ConsumerRecord<String, String> cr) throws Exception {
        messages.add(cr.value());
    }
    @GetMapping("/received")
    public String recv() throws Exception {
        return "Received " + messages.size() + " messages";
    }
}
```

# REST API

Easy way to produce and consume messages – limited performance

## # Get a list of topics

```
$ curl "http://kafkahost/topics"  
[{"name":"test","num_partitions":3}, {"name":"test2","num_partitions":1}]
```

## # Produce a message using JSON with the value '{ "foo": "bar" }' to the topic test

```
$ curl -X POST -H "Content-Type: application/vnd.kafka.json.v1+json" --data '{"records":[{"value":{"foo":"bar"} }]}' "http://kafkahost/topics/test"  
{"offsets":[{"partition":0,"offset":0,"error_code":null,"error":null}],"key_schema_id":null,"value_schema_id":null}
```

## # Create a consumer, read a message using the consumer, delete the consumer

```
$ curl -X POST -H "Content-Type: application/vnd.kafka.v1+json" --data '{"format": "json", "auto.offset.reset": "smallest"}' \  
http://kafkahost/consumers/my_json_consumer  
{"instance_id": "rest-consumer-11561681-8ba5-4b46-bed0-  
905ae1769bc6", "base_uri": "http://kafkahost/consumers/my_json_consumer/instances/rest-consumer-11561681-8ba5-4b46-bed0-  
905ae1769bc6"}  
  
$ curl -X GET -H "Accept: application/vnd.kafka.json.v1+json" http://kafkahost/consumers/my_json_consumer/instances/rest-consumer-  
11561681-8ba5-4b46-bed0-905ae1769bc6/topics/test  
[{"key":null,"value":{"foo":"bar"}, "partition":0, "offset":0}]  
  
$ curl -X DELETE \ http://localhost:8082/consumers/my_json_consumer/instances/rest-consumer-11561681-8ba5-4b46-bed0-905ae1769bc6
```

# IBM Event Streams – Admin REST API

API-based administrative controls for topics and bridges

GET /admin/topics

- returns the current list of topics

POST /admin/topics

- create a topic using values in request body

DELETE /admin/topics/{name}

- delete a topic with name {name}

GET /admin/bridges

- returns the current list of bridges

POST /admin/bridges

- creates a new instance of an Event Streams bridge

GET /admin/bridges/{name}

- get information about a bridge

DELETE /admin/bridges/{name}

- delete a bridge with name {name}

PATCH,PUT /admin/bridges/{name}

- update settings of bridge name {name}

# IBM Event Streams CLI

```
ibmcloud es init
ibmcloud es broker
ibmcloud es broker-config
ibmcloud es cluster
ibmcloud es topic
ibmcloud es topic-create
ibmcloud es topic-delete
ibmcloud es topic-delete-records
ibmcloud es topic-partitions-set
ibmcloud es topic-update
ibmcloud es topics
ibmcloud es group
ibmcloud es group-reset
ibmcloud es groups
ibmcloud es group-delete
```

# Lab Time